

# P3\_Cleaning\_Visualizations

January 31, 2026

Practical 3: Data Cleaning and Visualization

## 1 Data Cleaning and Visualization in Python

**Why this matters** - Real-world data is messy: missing values, wrong types, duplicates, outliers.  
- Cleaning is 70-80% of a data analyst's work. - Visualization helps explore patterns, detect issues, and communicate insights.

**Goals** 1. Load and explore a real dataset. 2. Clean common data problems systematically. 3. Create insightful visualizations.

### 1.0.1 1. Import the necessary libraries

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

#To suppress all warnings
import warnings
warnings.filterwarnings('ignore')
```

### 1.0.2 2. Loading and Exploration

```
[3]: # Load Titanic dataset
df = sns.load_dataset('titanic')

# have a quick look
print(df.shape)
df.head()
```

(891, 15)

```
[3]:   survived  pclass      sex    age  sibsp  parch      fare embarked  class \
0          0       3    male  22.0      1      0    7.2500      S  Third
1          1       1  female  38.0      1      0   71.2833      C  First
2          1       3  female  26.0      0      0    7.9250      S  Third
3          1       1  female  35.0      1      0   53.1000      S  First
4          0       3    male  35.0      0      0    8.0500      S  Third
```

```

      who  adult_male  deck  embark_town  alive  alone
0    man        True   NaN  Southampton    no  False
1  woman       False     C  Cherbourg   yes  False
2  woman       False   NaN  Southampton   yes   True
3  woman       False     C  Southampton   yes  False
4    man        True   NaN  Southampton    no   True

```

[5]: # Basic info  
df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   survived    891 non-null    int64  
 1   pclass      891 non-null    int64  
 2   sex         891 non-null    object  
 3   age         714 non-null    float64 
 4   sibsp       891 non-null    int64  
 5   parch       891 non-null    int64  
 6   fare         891 non-null    float64 
 7   embarked    889 non-null    object  
 8   class        891 non-null    category
 9   who          891 non-null    object  
 10  adult_male   891 non-null    bool   
 11  deck         203 non-null    category
 12  embark_town  889 non-null    object  
 13  alive        891 non-null    object  
 14  alone        891 non-null    bool  
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB

```

[7]: # Statistical summary (numeric only)  
df.describe()

	survived	pclass	age	sibsp	parch	fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
[9]: # Categorical summary  
df.describe(include=['object', 'category'])
```

```
[9]:      sex embarked class who deck embark_town alive  
count    891      889    891   891   203                 889    891  
unique     2        3      3     3     7                   3     2  
top      male       S  Third  man    C  Southampton    no  
freq     577      644    491   537    59                 644    549
```

```
[11]: # Check for missing values  
print("\nMissing values per column:")  
print(df.isnull().sum())
```

Missing values per column:

```
survived          0  
pclass            0  
sex               0  
age             177  
sibsp            0  
parch            0  
fare              0  
embarked         2  
class            0  
who               0  
adult_male       0  
deck             688  
embark_town      2  
alive             0  
alone             0  
dtype: int64
```

```
[13]: # check duplicates  
print(f"\nNumber of duplicate rows: {df.duplicated().sum()}")
```

Number of duplicate rows: 107

## 1.1 Data Cleaning

**Common issues in Titanic:** - Missing values: age (177), embarked (2), deck (688), embtown (not useful) - Data types: some could be categorical - Outliers - Redundant/unuseful columns

### 1.1.1 Handling Missing Values

```
[15]: # 1. Drop columns with too many missing values (>70%)  
df_clean = df.drop(columns=['deck'])  # 688 missing out of 891
```

```
[17]: # 2. Fill 'age' with median (robust to outliers)
median_age = df_clean['age'].median()
df_clean['age'].fillna(median_age, inplace=True)

[19]: # 3. Fill 'embarked' with mode (most frequent)
mode_embarked = df_clean['embarked'].mode()[0]
df_clean['embarked'].fillna(mode_embarked, inplace=True)

[21]: # 4. Fill 'embark_town' with mode
mode_embark_town = df_clean['embark_town'].mode()[0]
df_clean['embark_town'].fillna(mode_embark_town, inplace=True)

[23]: print("Missing values after cleaning:")
print(df_clean.isnull().sum())
```

Missing values after cleaning:

survived	0
pclass	0
sex	0
age	0
sibsp	0
parch	0
fare	0
embarked	0
class	0
who	0
adult_male	0
embark_town	0
alive	0
alone	0
dtype:	int64

### 1.1.2 Data Types and Categories

```
[26]: # Convert to categorical for memory and plotting
categorical_cols = ['sex', 'embarked', 'class', 'who', 'adult_male', 'alone']
for col in categorical_cols:
    df_clean[col] = df_clean[col].astype('category')
```

```
[28]: # Verify
df_clean.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          --          --      
 0   survived    891 non-null   int64
```

```

1  pclass          891 non-null      int64
2  sex             891 non-null      category
3  age             891 non-null      float64
4  sibsp           891 non-null      int64
5  parch           891 non-null      int64
6  fare             891 non-null      float64
7  embarked         891 non-null      category
8  class            891 non-null      category
9  who              891 non-null      category
10 adult_male       891 non-null      category
11 embark_town      891 non-null      object
12 alive            891 non-null      object
13 alone            891 non-null      category
dtypes: category(6), float64(2), int64(4), object(2)
memory usage: 61.7+ KB

```

### 1.1.3 Duplicates and Redundant Columns

```
[31]: # Remove exact duplicates (none in this dataset, but good practice)
df_clean.drop_duplicates(inplace=True)
```

```
[33]: # Drop redundant columns
df_clean.drop(columns=['alive', 'pclass', 'embark_town'], inplace=True, errors='ignore')
```

```
[35]: df_clean.head()
```

```

[35]:    survived     sex   age  sibsp  parch      fare embarked  class   who  \
0          0   male  22.0      1      0    7.2500      S  Third   man
1          1 female  38.0      1      0   71.2833      C  First  woman
2          1 female  26.0      0      0    7.9250      S  Third  woman
3          1 female  35.0      1      0   53.1000      S  First  woman
4          0   male  35.0      0      0    8.0500      S  Third   man

adult_male  alone
0      True  False
1     False  False
2     False  True
3     False  False
4      True  True

```

### 1.1.4 Outlier Detection

```
[38]: # Boxplot for fare and age
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
sns.boxplot(x=df_clean['fare'])
```


```

plt.title("Fare Distribution")

plt.subplot(1, 2, 2)
sns.boxplot(x=df_clean['age'])
plt.title("Age Distribution")

plt.show()

```



```

[40]: # Optional: Cap extreme fares using IQR method
Q1 = df_clean['fare'].quantile(0.25)
Q3 = df_clean['fare'].quantile(0.75)
IQR = Q3 - Q1
upper_bound = Q3 + 1.5 * IQR

print(f"Fares above {upper_bound:.2f} are potential outliers.")
df_clean['fare_capped'] = df_clean['fare'].clip(upper=upper_bound)

```

Fares above 73.42 are potential outliers.

## 1.2 Visualizations

We will create: - Distribution plots - Categorical counts - Relationships (scatter, box) - Correlation heatmap

### 1.2.1 Univariate Plots

```

[44]: plt.figure(figsize=(14, 4))

# Age distribution
plt.subplot(1, 3, 1)
sns.histplot(df_clean['age'], kde=True, bins=30)
plt.title('Age Distribution')

```


```

# Fare distribution (capped)
plt.subplot(1, 3, 2)
sns.histplot(df_clean['fare_capped'], kde=True, bins=30)
plt.title('Fare Distribution (Capped)')

# Survival rate
plt.subplot(1, 3, 3)
sns.countplot(x='survived', data=df_clean)
plt.title('Survival Count')

plt.tight_layout()
plt.show()

```



## 1.2.2 Categorical Relationships

```


[47]: plt.figure(figsize=(14, 5))

# Survival by sex
plt.subplot(1, 2, 1)
sns.countplot(x='sex', hue='survived', data=df_clean)
plt.title('Survival by Sex')

# Survival by class
plt.subplot(1, 2, 2)
sns.countplot(x='class', hue='survived', data=df_clean)
plt.title('Survival by Class')


plt.show()

```




### 1.2.3 Numerical Relationships

```
[50]: # Age vs Fare, colored by survival
sns.scatterplot(x='age', y='fare_capped', hue='survived', style='sex',
                 alpha=0.7, data=df_clean)
plt.title('Age vs Fare by Survival and Sex')
plt.show()
```




```
[52]: # Boxplot: Age by Class
sns.boxplot(x='class', y='age', hue='survived', data=df_clean)
plt.title('Age Distribution by Class and Survival')
plt.show()
```




#### 1.2.4 Pairplot and Heatmap

```
[55]: # Pairplot for key variables
sns.pairplot(df_clean, vars=['age', 'fare_capped', 'sibsp', 'parch'],
             hue='survived', diag_kind='kde')
plt.show()
```



```
[56]: # Correlation heatmap
plt.figure(figsize=(8, 6))
numeric_cols = df_clean.select_dtypes(include=np.number).columns
sns.heatmap(df_clean[numeric_cols].corr(), annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Heatmap')
plt.show()
```



[ ]: